



Intro to Deep Learning with Keras

Amol Mavuduru
AIS VP + Tech Lead





How this workshop is going to work...

- First - a very brief intro to machine learning and deep learning concepts.
- Second - a deep dive into deep learning theory and how to use Keras for deep learning.



Resources for this Workshop - on GitHub

<https://github.com/AmolMavuduru/IntroToDeepLearning>

Includes slides, code, and images used for examples.



Software Dependencies

You will need the following libraries to run the code:

- Anaconda distribution of Python
- Keras and TensorFlow.



+





About Keras and TensorFlow

TensorFlow

- Developed and released by the Google Brain Team in 2015.
- Arguably the most popular deep learning library.
- Dataflow programming paradigm (tough for beginners).



Keras

- Built by François Chollet, a Google AI researcher.
- Released in 2015.
- Object-oriented programming paradigm (easy for beginners).
- Can run on top of TensorFlow.





What is machine learning?

- Machine learning is an area of **artificial intelligence** that is focused on using **mathematical** and **statistical** techniques to “**give machines the ability to learn**” without being explicitly programmed.
- Usually involves **some form of data** with **examples** to learn from.




How Supervised ML Works

1. Start with training data
2. Train an algorithm on the training data.
3. Test the algorithm's performance on test data.
4. Repeat steps 2-3 with improved algorithms.

Supervised ML - Training Phase


- Training data has examples with **features** (inputs) and **targets** (outputs/what you want to predict).
- Example:
 - features - temperature during last 10 days
 - target - temperature for the next day

Training data



Supervised ML - Testing Phase


- Test data also has features and targets.
- Except this time, we are trying to predict the targets using the features and evaluate our predicts.





What is Deep Learning?

- A form of machine learning inspired by human learning.
- Uses algorithms called neural networks.
- Based on a lot of math!





Applications of Deep Learning

- Self-driving cars
- Cancer detection
- Image recognition
- Language translation
- Automatic text generation
- Speech recognition
- Stock market prediction
- Games (Chess, Go, etc.)
- Music composition



MNIST - Recognizing Handwritten Digits

- In this workshop we will see how to train a neural network to recognize handwritten digits!
- We will see a neural network reach human-level performance in just a few minutes of training!